

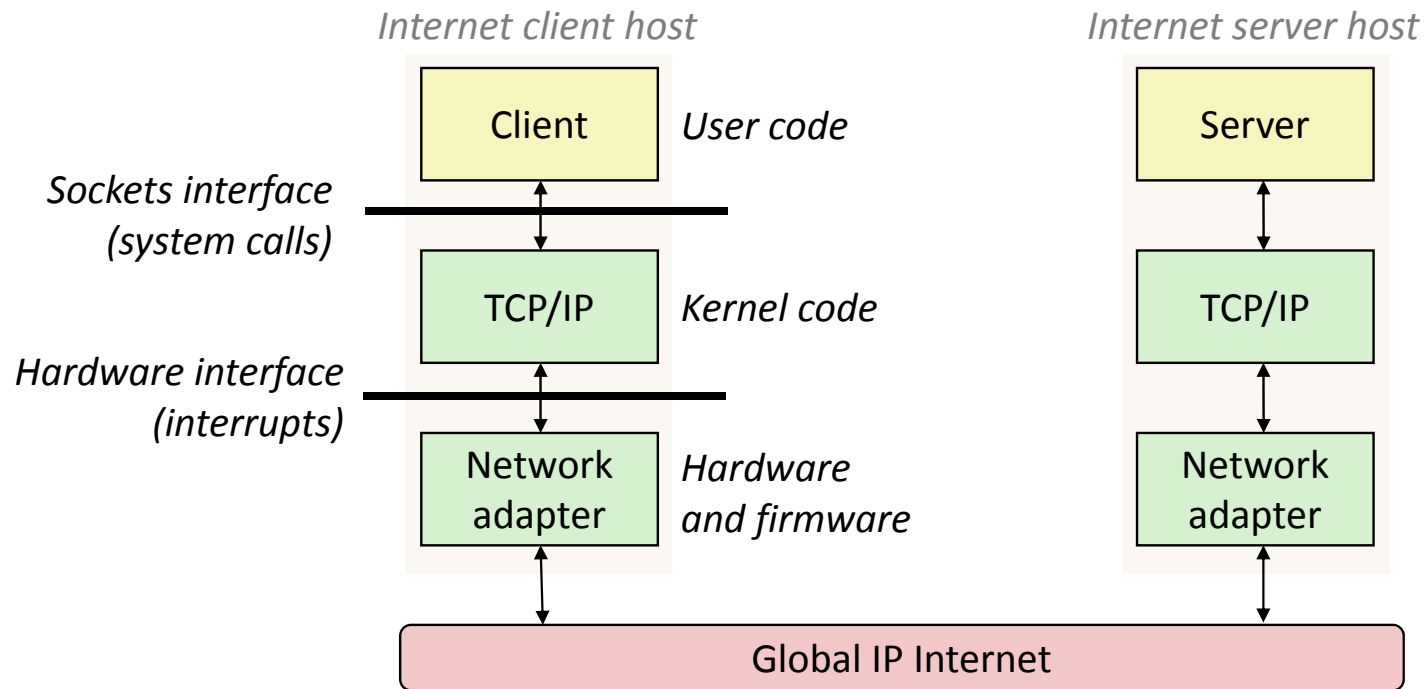
INTERNET

L4: Internet

# Last Lecture Re-cap: IP Internet

- Based on the TCP/IP protocol family
  - ▣ IP (Internet protocol) :
    - Provides *basic naming scheme (DNS)* and unreliable *delivery capability* of packets (datagrams) from host-to-host
  - ▣ UDP (Unreliable Datagram Protocol)
    - Uses IP to provide unreliable datagram delivery from *process-to-process*
  - ▣ TCP (Transmission Control Protocol)
    - Uses IP to provide *reliable* byte streams from process-to-process over connections
- Accessed via a mix of Unix file I/O and functions from the *sockets interface*

# Hardware and Software Organization of an Internet Application




# Today's Lecture

4

- TCP
  - ▣ Connection establishment, flow control, reliability, congestion control
- I/O
  - ▣ Unix I/O, (custom) RIO, standard I/O

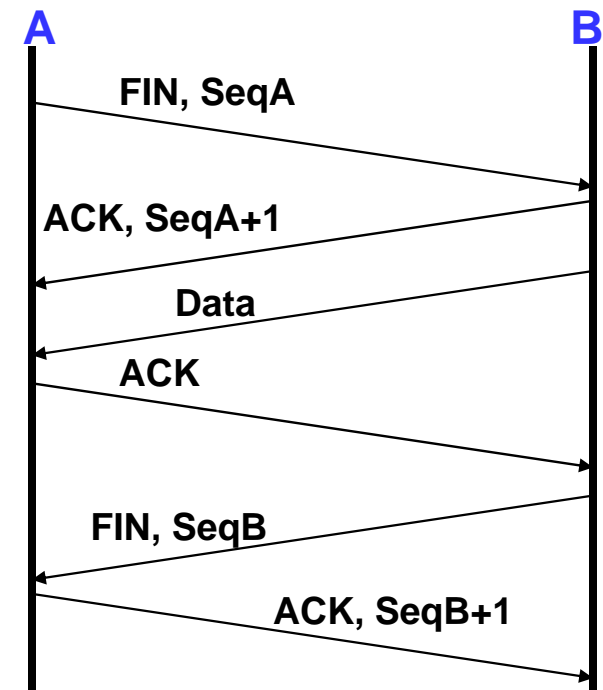
# Establishing Connection: Three-Way handshake

- Each side notifies other of starting sequence number it will use for sending
  - ▣ Why not simply chose 0?
    - Must avoid overlap with earlier incarnation
    - Security issues
- Each side acknowledges other's sequence number
  - ▣ SYN-ACK: Acknowledge sequence number + 1
- Can combine second SYN with first ACK

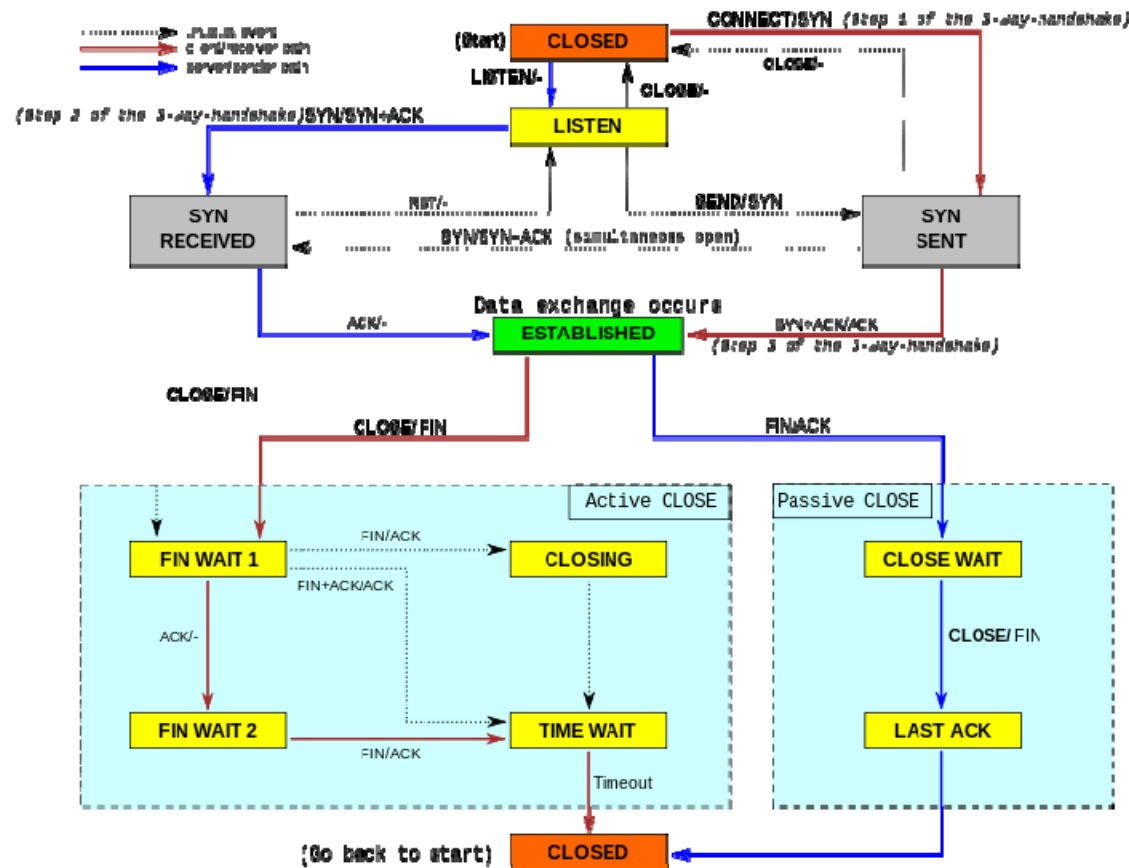


# Tearing Down Connection

- Either side can initiate tear down
  - ▣ Send FIN signal
  - ▣ “I’m not going to send any more data”
- Other side can continue sending data
  - ▣ Half open connection
  - ▣ Must continue to acknowledge
- Acknowledging FIN
  - ▣ Acknowledge last sequence number + 1

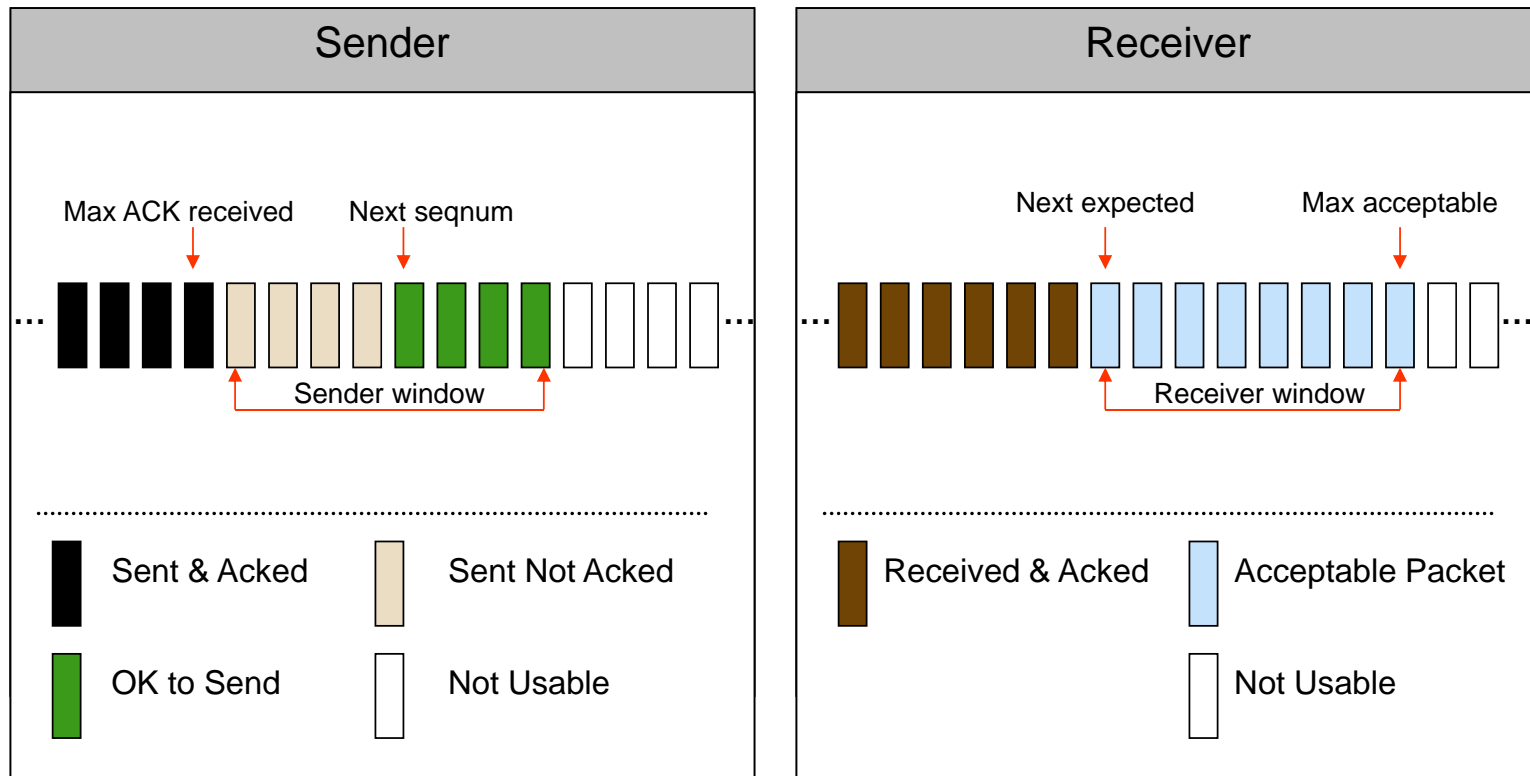


# TCP state transition (connection state)



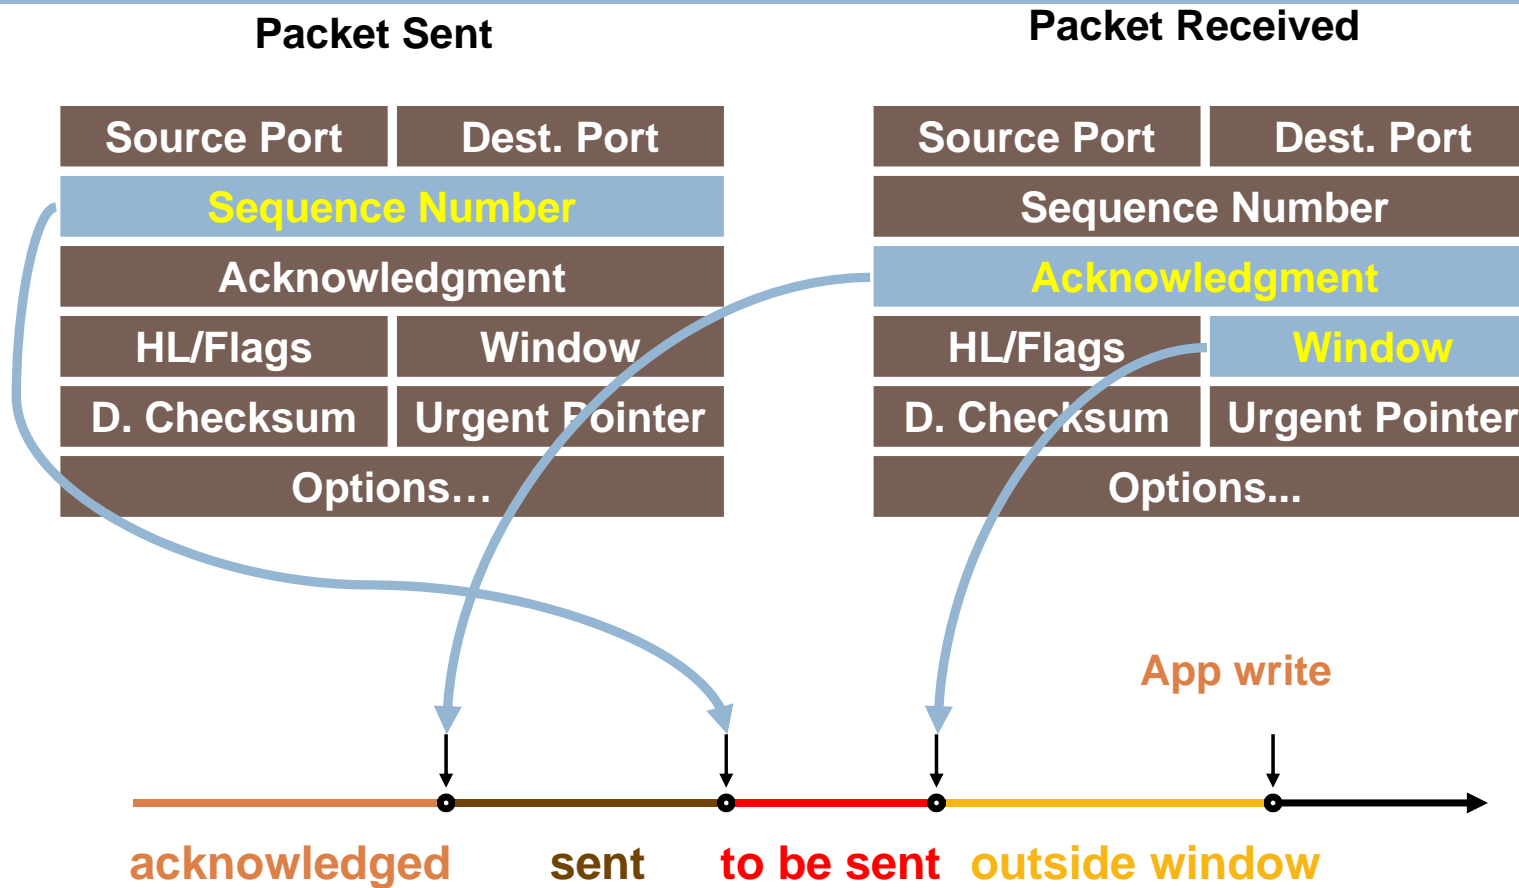
# Sender/Receiver State

8





# Window Flow Control: Send Side




# TCP congestion control

---

- Need to share network resources.
- But neither the sender or the receiver knows how much b/w is available.
- How much should we send?


# TCP congestion control

- Slow start
  - ▣ Increase the congestion window +1 for every ack.
- Fast recovery
  - ▣ On detection of dropped packet (dup ack) reduce the congestion window by half.
    - Called multiplicative decrease
- Congestion avoidance
  - ▣ Increase 1 every RTT
    - Called additive increase
- Details in computer networks course



# TCP Saw Tooth Behavior

12



# Important Lessons

- TCP state diagram → setup/teardown
  - ▣ Making sure both sides end up in same state
- TCP congestion control
  - ▣ Need to share some resources without knowing their current state
  - ▣ Good example of adapting to network performance
- Sliding window flow control
  - ▣ Addresses buffering issues and keeps link utilized
  - ▣ Need to ensure that distributed resources that are known about aren't overloaded

# Today

---

- TCP
- I/O
  - ▣ Unix I/O
  - ▣ RIO (robust I/O) package
  - ▣ Standard I/O

# Unix Files

- A Unix *file* is a sequence of  $m$  bytes:
  - $B_0, B_1, \dots, B_k, \dots, B_{m-1}$
- All I/O devices are represented as files:
  - `/dev/sda2` (`/usr` disk partition)
  - `/dev/tty2` (terminal)
- Even the kernel is represented as a file:
  - `/dev/kmem` (kernel memory image)
  - `/proc` (kernel data structures)

# Unix File Types

- Regular file
  - ▣ File containing user/app data (binary, text, whatever)
  - ▣ OS does not know anything about the format
    - other than “sequence of bytes”, akin to main memory
- Directory file
  - ▣ A file that contains the names and locations of other files
- Character special and block special files
  - ▣ Terminals (character special) and disks (block special)
- FIFO (named pipe)
  - ▣ A file type used for inter-process communication
- Socket
  - ▣ A file type used for network communication between processes



# Unix I/O

## □ Key Features

- Elegant mapping of files to devices allows kernel to export simple interface called Unix I/O
- Important idea: All input and output is handled in a consistent and uniform way

## □ Basic Unix I/O operations (system calls):

### □ Opening and closing files


- `open()` and `close()`

### □ Reading and writing a file

- `read()` and `write()`

### □ Changing the **current file position** (seek)

- indicates next offset into file to read or write
- `lseek()`



# Opening Files

- Opening a file informs the kernel that you are getting ready to access that file

```
int fd;    /* file descriptor */  
  
if ((fd = open("/etc/hosts", O_RDONLY)) < 0) {  
    perror("open");  
    exit(1);  
}
```

- Returns a small identifying integer *file descriptor*
  - ▣ `fd == -1` indicates that an error occurred
- Each process created by a Unix shell begins life with three open files associated with a terminal:
  - ▣ 0: standard input
  - ▣ 1: standard output
  - ▣ 2: standard error

# Closing Files

- Closing a file informs the kernel that you are finished accessing that file

```
int fd;      /* file descriptor */
int retval; /* return value */

if ((retval = close(fd)) < 0) {
    perror("close");
    exit(1);
}
```

- Closing an already closed file is a recipe for disaster in threaded programs (more on this later)
- Moral: Always check return codes, even for seemingly benign functions such as `close()`

# Reading Files

- Reading a file copies bytes from the current file position to memory, and then updates file position

```
char buf[512];
int fd;      /* file descriptor */
int nbytes;  /* number of bytes read */

/* Open file fd ... */
/* Then read up to 512 bytes from file fd */
if ((nbytes = read(fd, buf, sizeof(buf))) < 0) {
    perror("read");
    exit(1);
}
```

- Returns number of bytes read from file `fd` into `buf`
  - Return type `ssize_t` is signed integer
  - `nbytes < 0` indicates that an error occurred
  - **Short counts** (`nbytes < sizeof(buf)`) are possible and are not errors!

# Writing Files

- Writing a file copies bytes from memory to the current file position, and then updates current file position

```
char buf[512];
int fd;      /* file descriptor */
int nbytes;  /* number of bytes read */

/* Open the file fd ... */
/* Then write up to 512 bytes from buf to file fd */
if ((nbytes = write(fd, buf, sizeof(buf))) < 0) {
    perror("write");
    exit(1);
}
```

- Returns number of bytes written from buf to file fd
  - ▣ **nbytes** < 0 indicates that an error occurred
  - ▣ As with reads, short counts are possible and are not errors!

# Simple Unix I/O example

- Copying standard in to standard out, one byte at a time

```
#include "csapp.h"

int main(void)
{
    char c;

    while(Read(STDIN_FILENO, &c, 1) != 0)
        Write(STDOUT_FILENO, &c, 1);
    exit(0);
}
```

cpstdin.c

Note the use of error handling wrappers for read and write (Appendix A).

# Dealing with Short Counts

- Short counts can occur in these situations:
  - ▣ Encountering (end-of-file) EOF on reads
  - ▣ Reading text lines from a terminal
  - ▣ Reading and writing network sockets or Unix pipes
  
- Short counts never occur in these situations:
  - ▣ Reading from disk files (except for EOF)
  - ▣ Writing to disk files
  
- One way to deal with short counts in your code:
  - ▣ Use the RIO (Robust I/O) package from your textbook's `c_sapp.c` file (Appendix B)

# Today

---

- TCP
- I/O
  - ▣ Unix I/O
  - ▣ **RIO (robust I/O) package**
  - ▣ Standard I/O



# The RIO Package

- RIO is a set of wrappers that provide efficient and robust I/O in applications, such as network programs that are subject to short counts
- RIO provides two different kinds of functions
  - ▣ Unbuffered input and output of binary data
    - `rio_readn` and `rio_writen`
  - ▣ Buffered input of binary data and text lines
    - `rio_readlineb` and `rio_readnb`
    - Buffered RIO routines are thread-safe and can be interleaved arbitrarily on the same descriptor
- Download from <http://csapp.cs.cmu.edu/public/code.html>  
→ `src/csapp.c` and `include/csapp.h`

# Unbuffered RIO Input and Output

- Same interface as Unix `read` and `write`
- Especially useful for transferring data on network sockets

```
#include "csapp.h"
```

```
ssize_t rio_readn(int fd, void *usrbuf, size_t n);  
ssize_t rio_writen(int fd, void *usrbuf, size_t n);
```

Return: num. bytes transferred if OK, 0 on EOF (`rio_readn` only), -1 on error

- **`rio_readn`** returns short count only if it encounters EOF
  - Only use it when you know how many bytes to read
- **`rio_writen`** never returns a short count
- Calls to **`rio_readn`** and **`rio_writen`** can be interleaved arbitrarily on the same descriptor

# Implementation of `rio_readn`


```
/*
 * rio_readn - robustly read n bytes (unbuffered)
 */
ssize_t rio_readn(int fd, void *usrbuf, size_t n)
{
    size_t nleft = n;
    ssize_t nread;
    char *bufp = usrbuf;

    while (nleft > 0) {
        if ((nread = read(fd, bufp, nleft)) < 0) {
            if (errno == EINTR) /* interrupted by sig handler return */
                nread = 0;      /* and call read() again */
            else
                return -1;      /* errno set by read() */
        }
        else if (nread == 0)
            break;              /* EOF */
        nleft -= nread;
        bufp += nread;
    }
    return (n - nleft);         /* return >= 0 */
}
```

csapp.c


# Buffered I/O: Motivation

- Applications often read/write one character at a time
  - ▣ `getc`, `putc`, `ungetc`
  - ▣ `gets`, `fgets`
    - Read line of text one character at a time, stopping at newline
- Implementing as Unix I/O calls expensive
  - ▣ `read` and `write` require Unix kernel calls
    - > 10,000 clock cycles
- Solution: Buffered read
  - ▣ Use Unix `read` to grab block of bytes
  - ▣ User input functions take one byte at a time from buffer
    - Refill buffer when empty

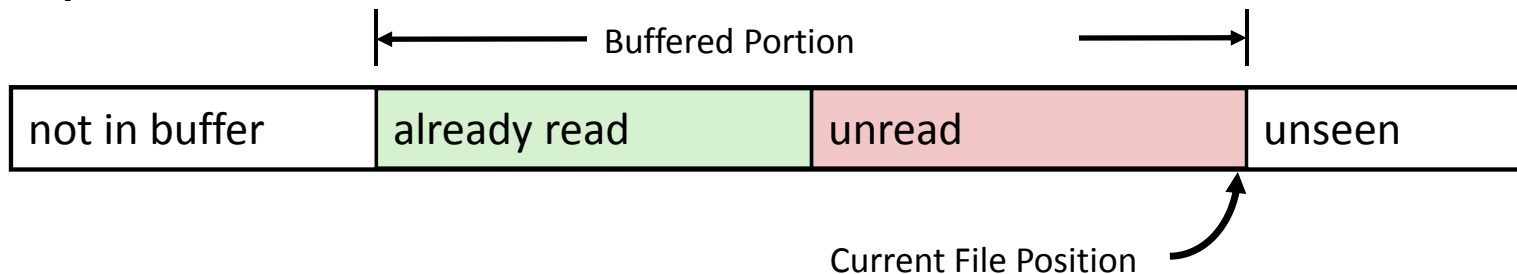


# Buffered I/O: Implementation

- For reading from file
- File has associated buffer to hold bytes that have been read from file but not yet read by user code




- Layered on Unix file:



# Buffered I/O: Declaration

- All information contained in struct



```
typedef struct {  
    int rio_fd;           /* descriptor for this internal buf */  
    int rio_cnt;          /* unread bytes in internal buf */  
    char *rio_bufptr;     /* next unread byte in internal buf */  
    char rio_buf[RIO_BUFSIZE]; /* internal buffer */  
} rio_t;
```

# Buffered RIO Input Functions

- Efficiently read text lines and binary data from a file partially cached in an internal memory buffer

```
#include "csapp.h"

void rio_readinitb(rio_t *rp, int fd);

ssize_t rio_readlineb(rio_t *rp, void *usrbuf, size_t maxlen);
```

Return: num. bytes read if OK, 0 on EOF, -1 on error

- ▣ **rio\_readlineb** reads a text line of up to **maxlen** bytes from file **fd** and stores the line in **usrbuf**
  - Especially useful for reading text lines from network sockets
- ▣ Stopping conditions
  - **maxlen** bytes read
  - EOF encountered
  - Newline ('\n') encountered

# Buffered RIO Input Functions (cont)

```
#include "csapp.h"
```

```
void rio_readinitb(rio_t *rp, int fd);
```

```
ssize_t rio_readlineb(rio_t *rp, void *usrbuf, size_t maxlen);
```

```
ssize_t rio_readnb(rio_t *rp, void *usrbuf, size_t n);
```

Return: num. bytes read if OK, 0 on EOF, -1 on error

- ▣ **rio\_readnb** reads up to **n** bytes from file **fd**
- ▣ Stopping conditions
  - **maxlen** bytes read
  - EOF encountered
- ▣ Calls to **rio\_readlineb** and **rio\_readnb** can be interleaved arbitrarily on the same descriptor
  - Warning: Don't interleave with calls to **rio\_readn**



# RIO Example

- Copying the lines of a text file from standard input to standard output

```
#include "csapp.h"

int main(int argc, char **argv)
{
    int n;
    rio_t rio;
    char buf[MAXLINE];

    Rio_readinitb(&rio, STDIN_FILENO);
    while((n = Rio_readlineb(&rio, buf, MAXLINE)) != 0)
        Rio_writen(STDOUT_FILENO, buf, n);
    exit(0);
}
```

cpfile.c

# Standard I/O Functions

- The C standard library (`libc.so`) contains a collection of higher-level *standard I/O* functions
  - ▣ Documented in Appendix B of K&R.
- Examples of standard I/O functions:
  - ▣ Opening and closing files (**`fopen`** and **`fclose`**)
  - ▣ Reading and writing bytes (**`fread`** and **`fwrite`**)
  - ▣ Reading and writing text lines (**`fgets`** and **`fputs`**)
  - ▣ Formatted reading and writing (**`fscanf`** and **`fprintf`**)

# Standard I/O Streams

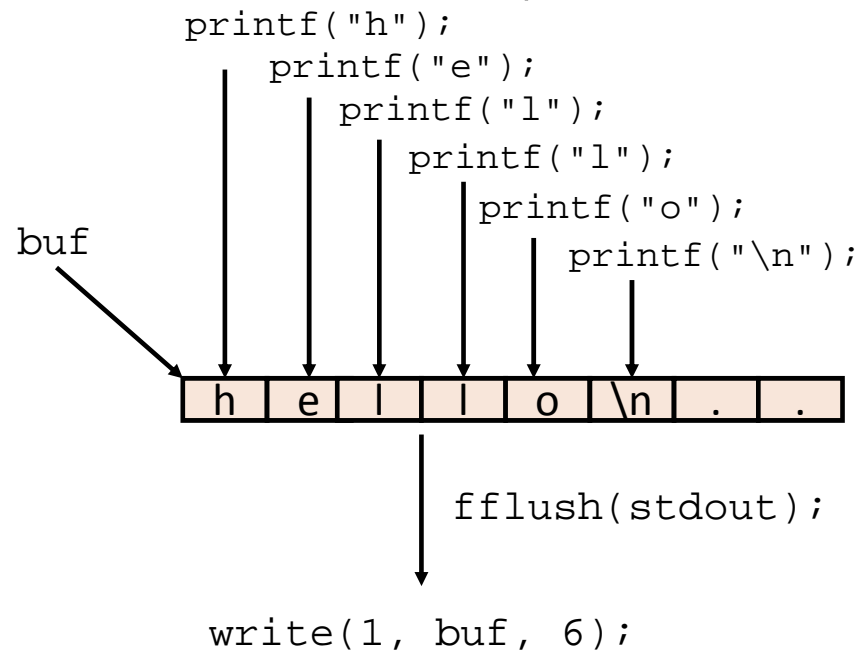
- Standard I/O models open files as *streams*
  - ▣ Abstraction for a file descriptor and a buffer in memory.
  - ▣ Similar to buffered RIO
- C programs begin life with three open streams (defined in `stdio.h`)
  - ▣ **`stdin`** (standard input)
  - ▣ **`stdout`** (standard output)
  - ▣ **`stderr`** (standard error)

```
#include <stdio.h>
extern FILE *stdin; /* standard input (descriptor 0) */
extern FILE *stdout; /* standard output (descriptor 1) */
extern FILE *stderr; /* standard error (descriptor 2) */

int main() {
    fprintf(stdout, "Hello, world\n");
}
```

# Buffering in Standard I/O

- Standard I/O functions use buffered I/O



- Buffer flushed to output fd on “\n” or fflush( ) call

# Standard I/O Buffering in Action

- You can see this buffering in action for yourself, using the always fascinating Unix `strace` program:

```
#include <stdio.h>

int main()
{
    printf("h");
    printf("e");
    printf("l");
    printf("l");
    printf("o");
    printf("\n");
    fflush(stdout);
    exit(0);
}
```

```
linux> strace ./hello
execve("./hello", ["hello"], [/* ... */]).
...
write(1, "hello\n", 6)                = 6
...
exit_group(0)                         = ?
```

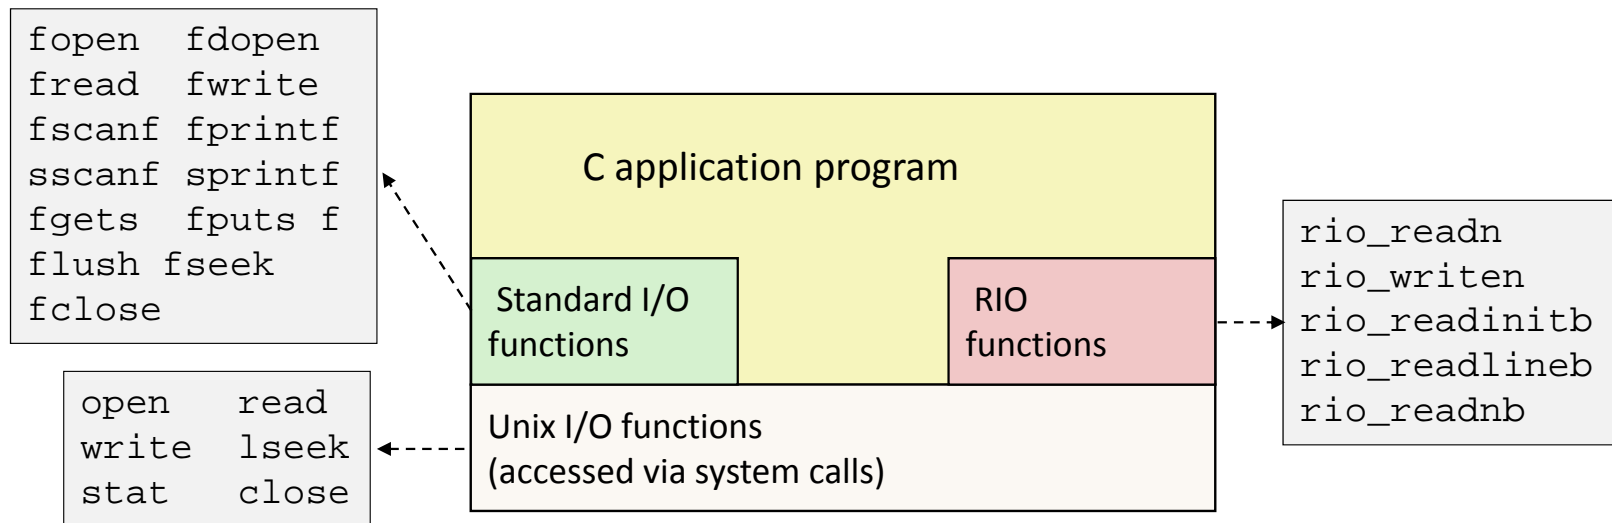
# Today

---

- TCP
- Unix I/O
- RIO (robust I/O) package
- Standard I/O
- **Conclusions**

# Unix I/O vs. Standard I/O vs. RIO

- Standard I/O and RIO are implemented using low-level Unix I/O



- Which ones should you use in your programs?

# Choosing I/O Functions

- General rule: use the highest-level I/O functions you can
  - ▣ Many C programmers are able to do all of their work using the standard I/O functions
- When to use standard I/O
  - ▣ When working with disk or terminal files
- When to use raw Unix I/O
  - ▣ Inside signal handlers, because Unix I/O is async-signal-safe.
  - ▣ In rare cases when you need absolute highest performance.
- When to use RIO
  - ▣ When you are reading and writing network sockets.
  - ▣ Avoid using standard I/O on sockets.



# A Programmer's View of the Internet

- Hosts are mapped to a set of 32-bit *IP addresses*
  - ▣ 128.2.203.179
- The set of IP addresses is mapped to a set of identifiers called Internet *domain names*
  - ▣ 128.2.203.179 is mapped to `www.cs.cmu.edu`
- A process on one Internet host can communicate with a process on another Internet host over a *connection*

# IP Addresses

- 32-bit IP addresses are stored in an *IP address struct*
  - ▣ IP addresses are always stored in memory in network byte order (big-endian byte order)
  - ▣ True in general for any integer transferred in a packet header from one machine to another.
    - E.g., the port number used to identify an Internet connection.

```
/* Internet address structure */
struct in_addr {
    unsigned int s_addr; /* network byte order (big-endian) */
};
```

Useful network byte-order conversion functions (“l” = 32 bits, “s” = 16 bits)

htonl: convert uint32\_t from host to network byte order  
htons: convert uint16\_t from host to network byte order  
ntohl: convert uint32\_t from network to host byte order  
ntohs: convert uint16\_t from network to host byte order

# Dotted Decimal Notation

- By convention, each byte in a 32-bit IP address is represented by its decimal value and separated by a period
  - IP address: `0x8002C2F2` = `128.2.194.242`
- Functions for converting between binary IP addresses and dotted decimal strings:
  - `inet_aton`: dotted decimal string → IP address in network byte order
  - `inet_ntoa`: IP address in network byte order → dotted decimal string
  - “n” denotes network representation
  - “a” denotes application representation

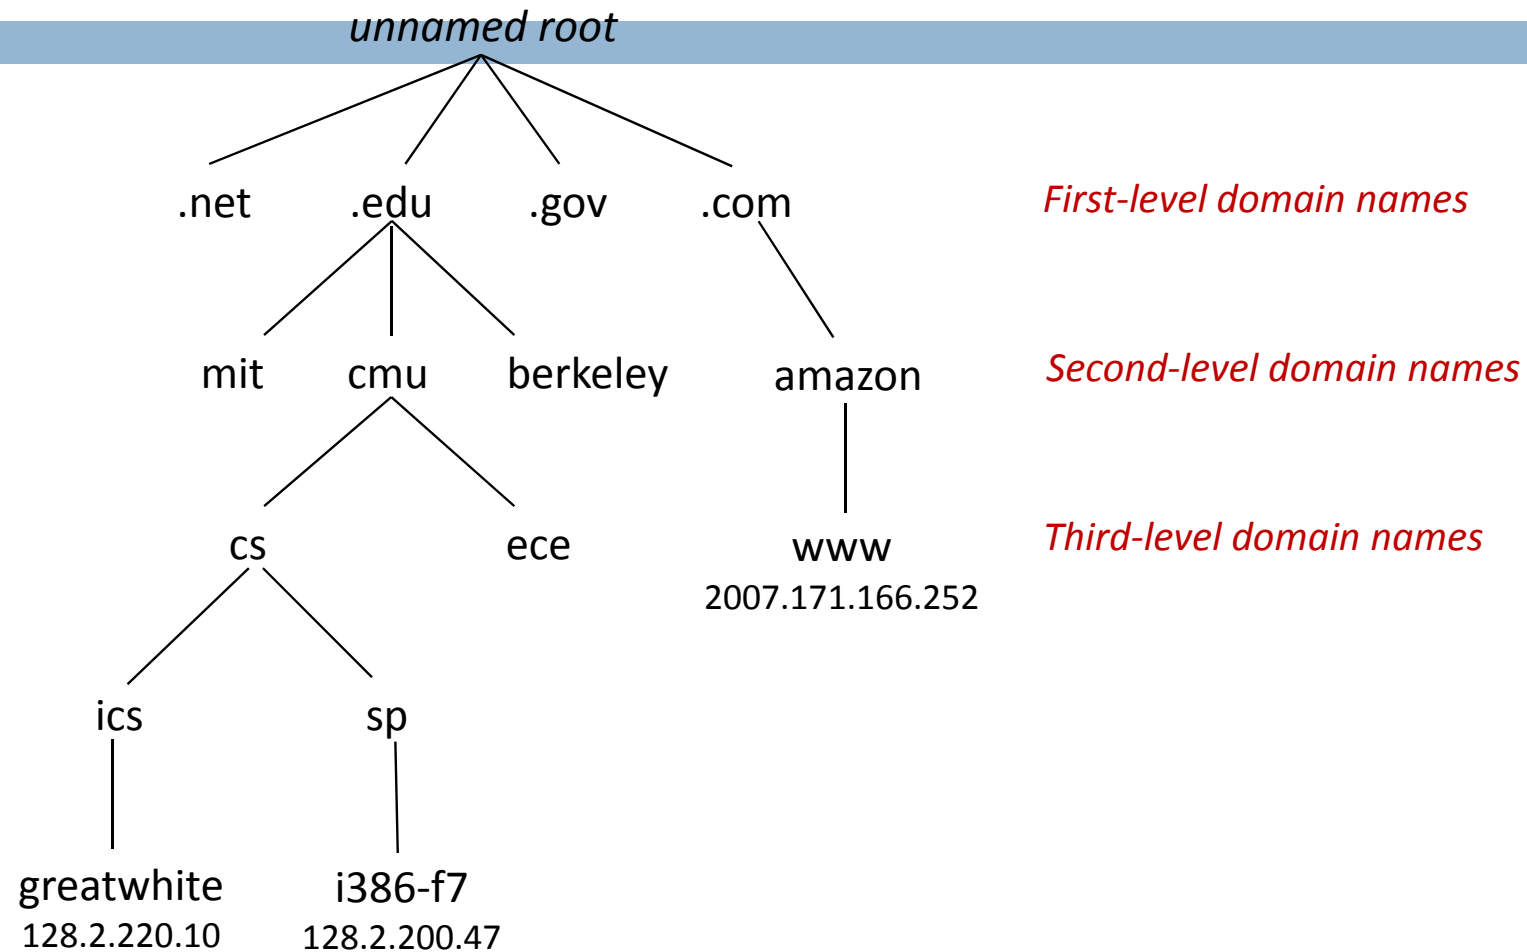
# IP Address Structure

- IP (V4) Address space divided into classes:

	0	1	2	3	8	16	24	31	
Class A	0	Net ID			Host ID				
Class B	1	0	Net ID				Host ID		
Class C	1	1	0	Net ID				Host ID	
Class D	1	1	1	0	Multicast address				
Class E	1	1	1	1	Reserved for experiments				

- Network ID Written in form w.x.y.z/n
  - ▣ n = number of bits in host address
  - ▣ E.g., CMU written as 128.2.0.0/16
    - Class B address
- Unrouted (private) IP addresses:  
10.0.0.0/8   172.16.0.0/12   192.168.0.0/16

# Internet Domain Names



# Domain Naming System (DNS)

- The Internet maintains a mapping between IP addresses and domain names in a huge worldwide distributed database called *DNS*
- ▣ Conceptually, programmers can view the DNS database as a collection of millions of *host entry structures*:

```
/* DNS host entry structure */
struct hostent {
    char    *h_name;           /* official domain name of host */
    char    **h_aliases;       /* null-terminated array of domain names */
    int     h_addrtype;        /* host address type (AF_INET) */
    int     h_length;          /* length of an address, in bytes */
    char    **h_addr_list;     /* null-terminated array of in_addr structs */
};
```

- Functions for retrieving host entries from DNS:
  - ▣ **gethostbyname**: query key is a DNS domain name.
  - ▣ **gethostbyaddr**: query key is an IP address.

# Properties of DNS Host Entries

- Each host entry is an equivalence class of domain names and IP addresses
- Each host has a locally defined domain name `localhost` which always maps to the *loopback address* `127.0.0.1`
- Different kinds of mappings are possible:
  - ▣ Simple case: one-to-one mapping between domain name and IP address:
    - `greatwhile.ics.cs.cmu.edu` maps to `128.2.220.10`
  - ▣ Multiple domain names mapped to the same IP address:
    - `eeecs.mit.edu` and `cs.mit.edu` both map to `18.62.1.6`
  - ▣ Multiple domain names mapped to multiple IP addresses:
    - `google.com` maps to multiple IP addresses
  - ▣ Some valid domain names don't map to any IP address:
    - for example: `ics.cs.cmu.edu`

# A Program That Queries DNS

```
int main(int argc, char **argv) { /* argv[1] is a domain name */
    char **pp;                    /* or dotted decimal IP addr */
    struct in_addr addr;
    struct hostent *hostp;

    if (inet_aton(argv[1], &addr) != 0)
        hostp = Gethostbyaddr((const char *)&addr, sizeof(addr),
                               AF_INET);
    else
        hostp = Gethostbyname(argv[1]);
    printf("official hostname: %s\n", hostp->h_name);

    for (pp = hostp->h_aliases; *pp != NULL; pp++)
        printf("alias: %s\n", *pp);

    for (pp = hostp->h_addr_list; *pp != NULL; pp++) {
        addr.s_addr = ((struct in_addr *)*pp)->s_addr;
        printf("address: %s\n", inet_ntoa(addr));
    }
}
```



# Using DNS Program

```
linux> ./dns greatwhite.ics.cs.cmu.edu  
official hostname: greatwhite.ics.cs.cmu.edu  
address 128.2.220.10
```

```
linux> ./dns 128.2.220.11  
official hostname: ANGELSHARK.ICS.CS.CMU.EDU  
address: 128.2.220.11
```

```
linux> ./dns www.google.com  
official hostname: www.l.google.com  
alias: www.google.com  
address: 72.14.204.99  
address: 72.14.204.103  
address: 72.14.204.104  
address: 72.14.204.147  
linux> dig +short -x 72.14.204.103  
iad04s01-in-f103.1e100.net.
```

# Querying DIG

- Domain Information Groper (dig) provides a scriptable command line interface to DNS

```
linux> dig +short greatwhite.ics.cs.cmu.edu
128.2.220.10
linux> dig +short -x 128.2.220.11
ANGELSHARK.ICS.CS.CMU.EDU.
linux> dig +short google.com
72.14.204.104
72.14.204.147
72.14.204.99
72.14.204.103
linux> dig +short -x 72.14.204.103
iad04s01-in-f103.1e100.net.
```

# More Exotic Features of DIG

- Provides more information than you would ever want about DNS

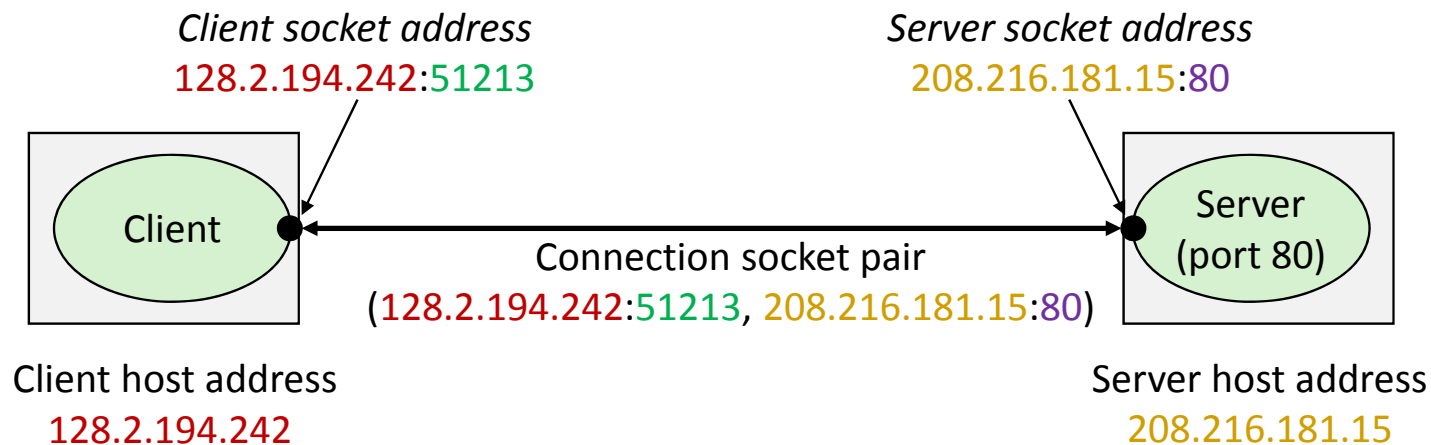
```
linux> dig www.phys.msu.ru a +trace  
128.2.220.10
```

```
linux> dig www.google.com a +trace
```

# Internet Connections

- Clients and servers communicate by sending streams of bytes over **connections**:
  - ▣ Point-to-point, full-duplex (2-way communication), and reliable.
- A **socket** is an endpoint of a connection
  - ▣ Socket address is an **IPAddress:port** pair
- A **port** is a 16-bit integer that identifies a process:
  - ▣ **Ephemeral port**: Assigned automatically on client when client makes a connection request
  - ▣ **Well-known port**: Associated with some service provided by a server (e.g., port 80 is associated with Web servers)
- A connection is uniquely identified by the socket addresses of its endpoints (**socket pair**)
  - ▣ **(cliaddr:cliport, servaddr:servport)**

# Putting it all Together: Anatomy of an Internet Connection



# Servers

- Servers are long-running processes (daemons)
  - ▣ Created at boot-time (typically) by the init process (process 1)
  - ▣ Run continuously until the machine is turned off
- Each server waits for requests to arrive on a well-known port associated with a particular service
  - ▣ Port 7: echo server
  - ▣ Port 23: telnet server
  - ▣ Port 25: mail server
  - ▣ Port 80: HTTP server
- A machine that runs a server process is also often referred to as a “server”

# Server Examples

- Web server (port 80)
  - Resource: files/compute cycles (CGI programs)
  - Service: retrieves files and runs CGI programs on behalf of the client
  
- FTP server (20, 21)
  - Resource: files
  - Service: stores and retrieve files
  
- Telnet server (23)
  - Resource: terminal
  - Service: proxies a terminal on the server machine
  
- Mail server (25)
  - Resource: email “spool” file
  - Service: stores mail messages in spool file

See `/etc/services` for a comprehensive list of the port mappings on a Linux machine

# Sockets Interface

- Created in the early 80's as part of the original Berkeley distribution of Unix that contained an early version of the Internet protocols
- Provides a user-level interface to the network
- Underlying basis for all Internet applications
- Based on client/server programming model




# Sockets

- What is a socket?

- To the kernel, a socket is an endpoint of communication
- To an application, a socket is a file descriptor that lets the application read/write from/to the network
  - **Remember:** All Unix I/O devices, including networks, are modeled as files

- Clients and servers communicate with each other by reading from and writing to socket descriptors



- The main distinction between regular file I/O and socket I/O is how the application “opens” the socket descriptors

# Example: Echo Client and Server

On Client

On Server

```
greatwhite> ./echoserveri 15213
```

```
linux> echoclient greatwhite.ics.cs.cmu.edu 15213
```

```
server connected to BRYANT-TP4.VLSI.CS.CMU.EDU (128.2.  
213.29), port 64690
```

```
type: hello there
```

```
server received 12 bytes
```

```
echo: HELLO THERE  
type: ^D
```

```
Connection closed
```

# Watching Echo Client / Server



Capturing from Microsoft - Wireshark

File Edit View Go Capture Analyze Statistics Telephony Tools Help

Filter: tcp.port eq 15213 Expression... Clear Apply

No.	Time	Source	Destination	Protocol	Info
1255	15.881493	128.237.252.163	128.2.220.10	TCP	55306 > 15213 [SYN] Seq=0 win=65535 Len=0 MSS=1
1256	15.883817	128.2.220.10	128.237.252.163	TCP	15213 > 55306 [SYN, ACK] Seq=0 Ack=1 win=5840 L
1257	15.883897	128.237.252.163	128.2.220.10	TCP	55306 > 15213 [ACK] Seq=1 Ack=1 win=65532 Len=0
1799	21.914380	128.237.252.163	128.2.220.10	TCP	55306 > 15213 [PSH, ACK] Seq=1 Ack=1 win=65532
1800	21.916474	128.2.220.10	128.237.252.163	TCP	15213 > 55306 [ACK] Seq=1 Ack=19 win=5888 Len=0
1801	21.916534	128.2.220.10	128.237.252.163	TCP	15213 > 55306 [PSH, ACK] Seq=1 Ack=19 win=5888
1816	22.112223	128.237.252.163	128.2.220.10	TCP	55306 > 15213 [ACK] Seq=19 Ack=19 win=65516 Len
2301	29.053184	128.237.252.163	128.2.220.10	TCP	55306 > 15213 [PSH, ACK] Seq=19 Ack=19 win=6551
2302	29.055004	128.2.220.10	128.237.252.163	TCP	15213 > 55306 [PSH, ACK] Seq=19 Ack=43 win=5888
2316	29.253626	128.237.252.163	128.2.220.10	TCP	55306 > 15213 [ACK] Seq=43 Ack=43 win=65492 Len
2382	30.229193	128.237.252.163	128.2.220.10	TCP	55306 > 15213 [FIN, ACK] Seq=43 Ack=43 win=6549

Frame 1799: 72 bytes on wire (576 bits), 72 bytes captured (576 bits)

Ethernet II, Src: Intel\_e3:54:e6 (00:16:ea:e3:54:e6), Dst: Carnegie\_20:00:64 (08:00:7f:20:00:64)

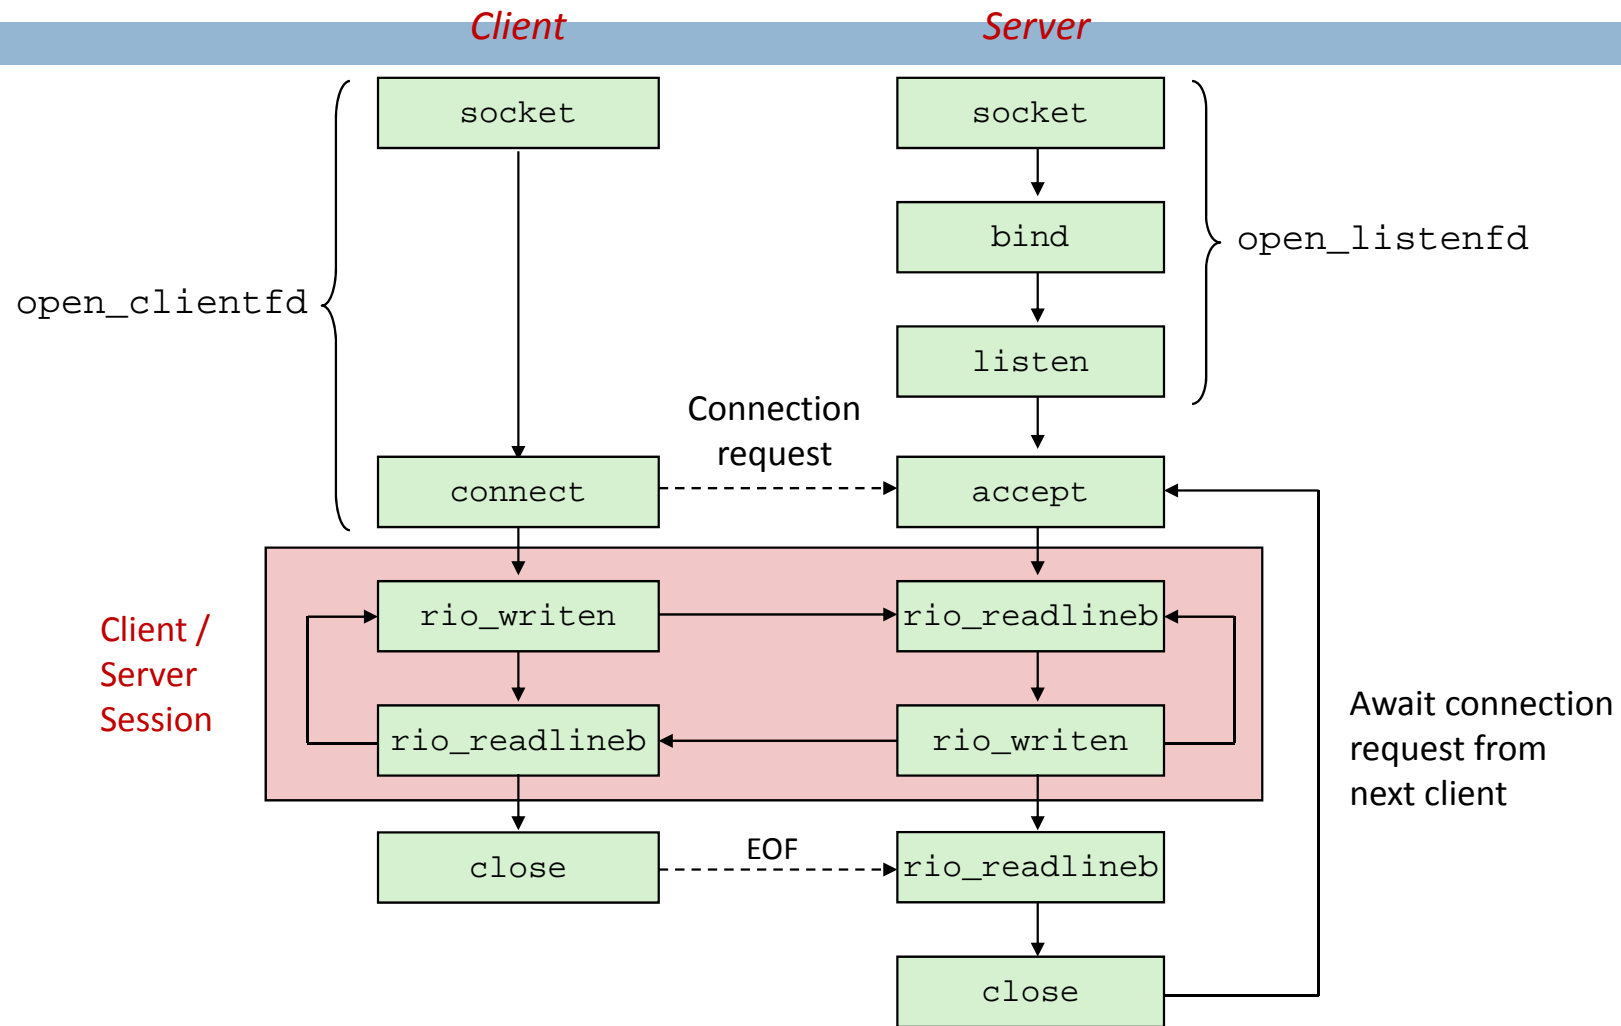
Internet Protocol, Src: 128.237.252.163 (128.237.252.163), Dst: 128.2.220.10 (128.2.220.10)

Transmission Control Protocol, Src Port: 55306 (55306), Dst Port: 15213 (15213), Seq: 1, Ack: 1, Len: 18

```
0000 08 00 7f 20 00 64 00 16 ea e3 54 e6 08 00 45 00 ... .d.. ..T...E.
0010 00 3a 2c 7a 40 00 80 06 f4 a5 80 ed fc a3 80 02 .:,z@... .....
0020 dc 0a d8 0a 3b 6d f4 a4 99 6c 75 de 71 6a 50 18 ....;m.. .lu.qjP.
0030 3f ff 96 8b 00 00 68 65 72 65 20 69 73 20 61 20 ?.....he re is a
0040 6d 65 73 73 61 67 65 0a message.
```

Microsoft: <live capture in progress> File: C:... Packets: 6950 Displayed: 13 Marked: 0 Profile: Default

# Overview of the Sockets Interface




# Socket Address Structures

## □ Generic socket address:

- For address arguments to **connect**, **bind**, and **accept**
- Necessary only because C did not have generic **(void \*)** pointers when the sockets interface was designed

```
struct sockaddr {  
    unsigned short  sa_family;    /* protocol family */  
    char            sa_data[14]; /* address data.  */  
};
```

sa\_family



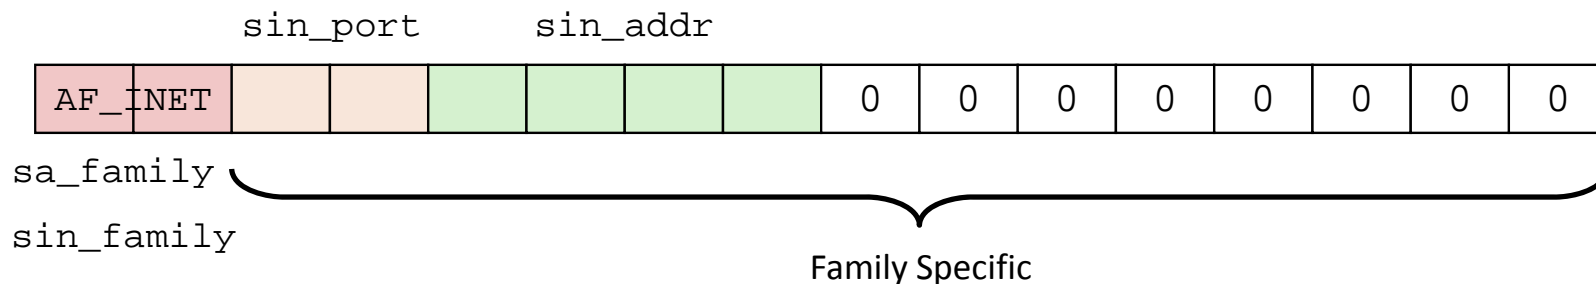
Family Specific

# Socket Address Structures

## Internet-specific socket address:

- Must cast (`sockaddr_in *`) to (`sockaddr *`) for `connect`, `bind`, and `accept`

```
struct sockaddr_in {  
    unsigned short  sin_family; /* address family (always AF_INET) */  
    unsigned short  sin_port;   /* port num in network byte order */  
    struct in_addr  sin_addr;   /* IP addr in network byte order */  
    unsigned char   sin_zero[8]; /* pad to sizeof(struct sockaddr) */  
};
```



# Echo Client Main Routine

```
#include "csapp.h"

/* usage: ./echoclient host port */
int main(int argc, char **argv)
{
    int clientfd, port;
    char *host, buf[MAXLINE];
    rio_t rio;
    host = argv[1]; port = atoi(argv[2]);
    clientfd = Open_clientfd(host, port);
    Rio_readinitb(&rio, clientfd);
    printf("type:"); fflush(stdout);
    while (Fgets(buf, MAXLINE, stdin) != NULL) {
        Rio_writen(clientfd, buf, strlen(buf));
        Rio_readlineb(&rio, buf, MAXLINE);
        printf("echo:");
        Fputs(buf, stdout);
        printf("type:"); fflush(stdout);
    }
    Close(clientfd);
    exit(0);
}
```

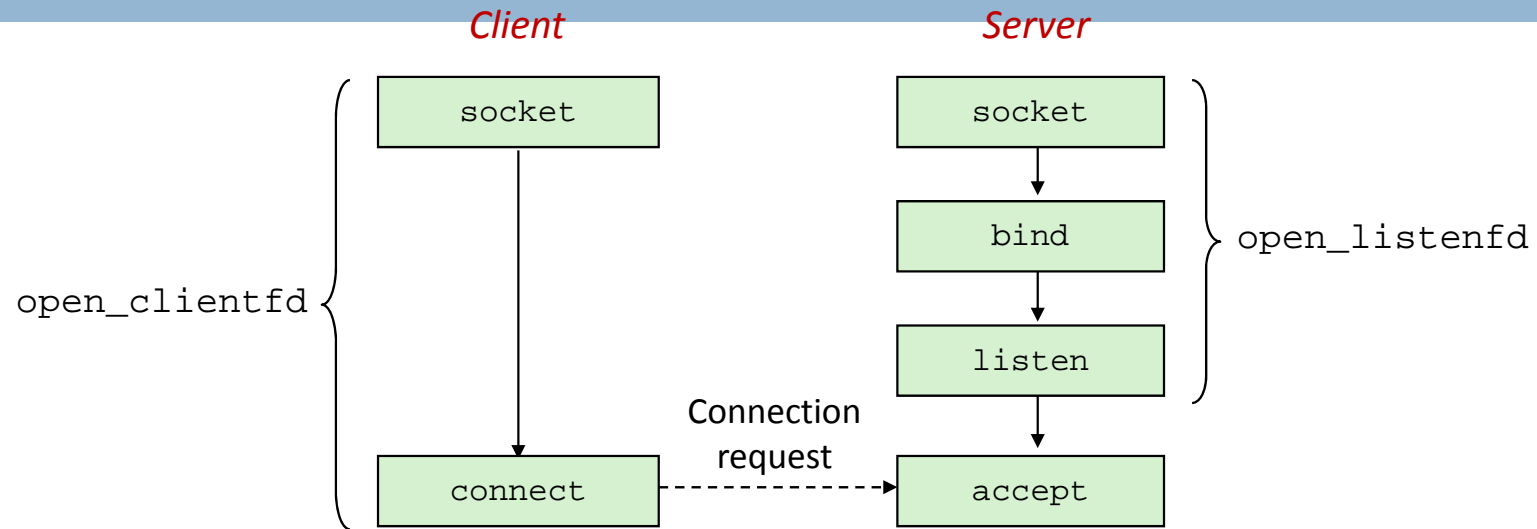
Send line to server

Receive line from server

Read input line

Print server response

# Overview of the Sockets Interface





# Echo Client: open\_clientfd

```
int open_clientfd(char *hostname, int port) {
    int clientfd;
    struct hostent *hp;
    struct sockaddr_in serveraddr;

    if ((clientfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)
        return -1; /* check errno for cause of error */

    /* Fill in the server's IP address and port */
    if ((hp = gethostbyname(hostname)) == NULL)
        return -2; /* check h_errno for cause of error */
    bzero((char *) &serveraddr, sizeof(serveraddr));
    serveraddr.sin_family = AF_INET;
    bcopy((char *)hp->h_addr_list[0],
          (char *)&serveraddr.sin_addr.s_addr, hp->h_length);
    serveraddr.sin_port = htons(port);

    /* Establish a connection with the server */
    if (connect(clientfd, (SA *) &serveraddr,
                sizeof(serveraddr)) < 0)
        return -1;
    return clientfd;
}
```

This function opens a connection from the client to the server at hostname:port

} Create socket

} Create address

} Establish connection

# Echo Client: `open_clientfd` (`socket`)

- `socket` creates a socket descriptor on the client
  - Just allocates & initializes some internal data structures
  - `AF_INET`: indicates that the socket is associated with Internet protocols
  - `SOCK_STREAM`: selects a reliable byte stream connection
    - provided by TCP

```
int clientfd; /* socket descriptor */

if ((clientfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)
    return -1; /* check errno for cause of error */

... <more>
```

# Echo Client: `open_clientfd` (`gethostbyname`)


- The client then builds the server's Internet address

```
int clientfd;                /* socket descriptor */
struct hostent *hp;          /* DNS host entry */
struct sockaddr_in serveraddr; /* server's IP address */

...

/* fill in the server's IP address and port */
if ((hp = gethostbyname(hostname)) == NULL)
    return -2; /* check h_errno for cause of error */
bzero((char *) &serveraddr, sizeof(serveraddr));
serveraddr.sin_family = AF_INET;
serveraddr.sin_port = htons(port);
bcopy((char *)hp->h_addr_list[0],
      (char *)&serveraddr.sin_addr.s_addr, hp->h_length);
```

Check  
this out!



# A Careful Look at bcopy Argument

```
/* DNS host entry structure */
struct hostent {
    . . .
    int      h_length;      /* length of an address, in bytes */
    char     **h_addr_list; /* null-terminated array of in_addr structs */
};
```

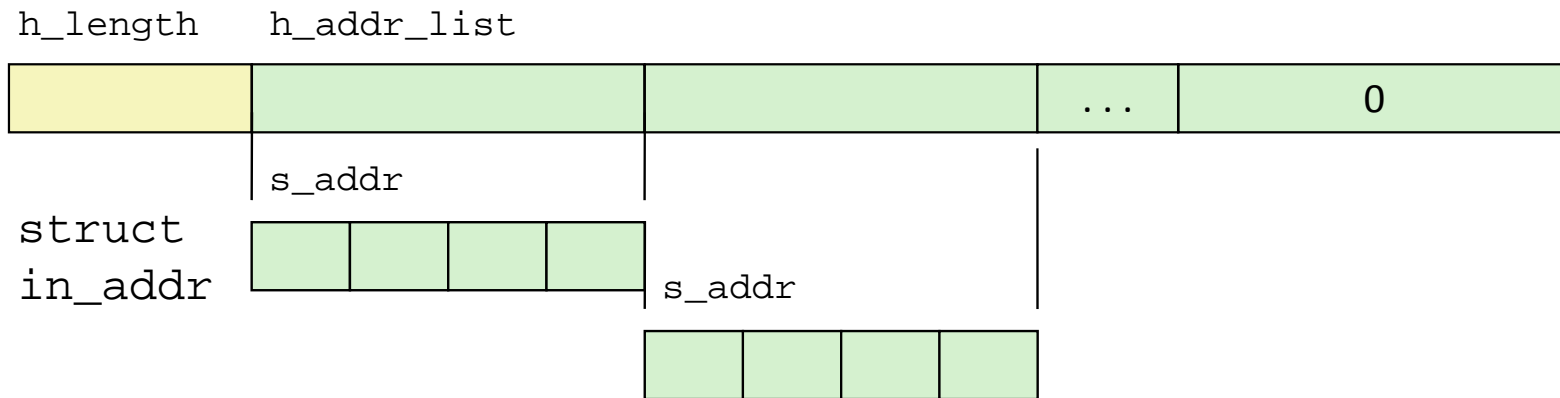
```
struct sockaddr_in {
    . . .
    struct in_addr  sin_addr; /* IP addr in network byte order */
    . . .
};
```

```
/* Internet address structure */
struct in_addr {
    unsigned int s_addr; /* network byte order (big-endian) */
};
```


```
struct hostent *hp; /* DNS host entry */
struct sockaddr_in serveraddr; /* server's IP address */
...
bcopy((char *)hp->h_addr_list[0], /* src, dest */
      (char *)&serveraddr.sin_addr.s_addr, hp->h_length);
```

# Bcopy Argument Data Structures

struct hostent



struct sockaddr\_in



# Echo Client: `open_clientfd`

(`connect`)

- Finally the client creates a connection with the server
  - ▣ Client process suspends (blocks) until the connection is created
  - ▣ After resuming, the client is ready to begin exchanging messages with the server via Unix I/O calls on descriptor `clientfd`

```
int clientfd;                /* socket descriptor */
struct sockaddr_in serveraddr; /* server address */
typedef struct sockaddr SA;    /* generic sockaddr */
...
/* Establish a connection with the server */
if (connect(clientfd, (SA *)&serveraddr, sizeof(serveraddr)) < 0)
    return -1;
return clientfd;
}
```

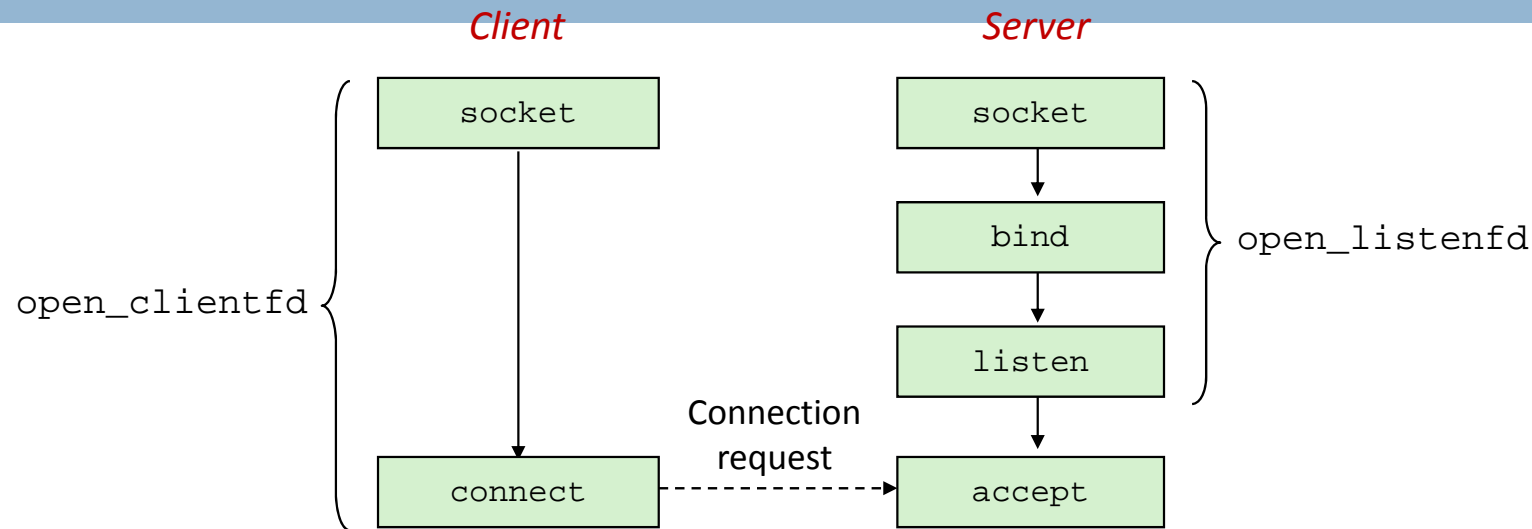
# Echo Server: Main Routine

```
int main(int argc, char **argv) {
    int listenfd, connfd, port, clientlen;
    struct sockaddr_in clientaddr;
    struct hostent *hp;
    char *haddrp;
    unsigned short client_port;

    port = atoi(argv[1]); /* the server listens on a port passed
                           on the command line */
    listenfd = open_listenfd(port);

    while (1) {
        clientlen = sizeof(clientaddr);
        connfd = Accept(listenfd, (SA *)&clientaddr, &clientlen);
        hp = Gethostbyaddr((const char *)&clientaddr.sin_addr.s_addr,
                           sizeof(clientaddr.sin_addr.s_addr), AF_INET);
        haddrp = inet_ntoa(clientaddr.sin_addr);
        client_port = ntohs(clientaddr.sin_port);
        printf("server connected to %s (%s), port %u\n",
               hp->h_name, haddrp, client_port);
        echo(connfd);
        Close(connfd);
    }
}
```

# Overview of the Sockets Interface



## □ Office Telephone Analogy for Server

- Socket: Buy a phone
- Bind: Tell the local administrator what number you want to use
- Listen: Plug the phone in
- Accept: Answer the phone when it rings



# Echo Server: open\_listenfd

```
int open_listenfd(int port)
{
    int listenfd, optval=1;
    struct sockaddr_in serveraddr;

    /* Create a socket descriptor */
    if ((listenfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)
        return -1;

    /* Eliminates "Address already in use" error from bind. */
    if (setsockopt(listenfd, SOL_SOCKET, SO_REUSEADDR,
                    (const void *)&optval , sizeof(int)) < 0)
        return -1;

    ... <more>
```

## Echo Server: open\_listenfd (cont.)

...

```
/* Listenfd will be an endpoint for all requests to port
   on any IP address for this host */
bzero((char *) &serveraddr, sizeof(serveraddr));
serveraddr.sin_family = AF_INET;
serveraddr.sin_addr.s_addr = htonl(INADDR_ANY);
serveraddr.sin_port = htons((unsigned short)port);
if (bind(listenfd, (SA *)&serveraddr, sizeof(serveraddr)) < 0)
    return -1;

/* Make it a listening socket ready to accept
   connection requests */
if (listen(listenfd, LISTENQ) < 0)
    return -1;

return listenfd;
}
```

# Echo Server: `open_listenfd` (`socket`)

- `socket` creates a socket descriptor on the server
  - ▣ **`AF_INET`**: indicates that the socket is associated with Internet protocols
  - ▣ **`SOCK_STREAM`**: selects a reliable byte stream connection (TCP)

```
int listenfd; /* listening socket descriptor */

/* Create a socket descriptor */
if ((listenfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)
    return -1;
```

# Echo Server: `open_listenfd` (`setsockopt`)

- The socket can be given some attributes

```
...  
/* Eliminates "Address already in use" error from bind(). */  
if (setsockopt(listenfd, SOL_SOCKET, SO_REUSEADDR,  
              (const void *)&optval , sizeof(int)) < 0)  
    return -1;
```


- Handy trick that allows us to rerun the server immediately after we kill it
  - ▣ Otherwise we would have to wait about 15 seconds
  - ▣ Eliminates “Address already in use” error from `bind()`
- Strongly suggest you do this for all your servers to simplify debugging

# Echo Server: `open_listenfd`

## (initialize socket address)

- Initialize socket with server port number
- Accept connection from any IP address

```
struct sockaddr_in serveraddr; /* server's socket addr */
...
/* listenfd will be an endpoint for all requests to port
   on any IP address for this host */
bzero((char *) &serveraddr, sizeof(serveraddr));
serveraddr.sin_family = AF_INET;
serveraddr.sin_port = htons((unsigned short)port);
serveraddr.sin_addr.s_addr = htonl(INADDR_ANY);
```



sa\_family

sin\_family

# Echo Server: open\_listenfd

(bind)

- bind associates the socket with the socket address we just created

```
int listenfd; /* listening socket */
struct sockaddr_in serveraddr; /* server's socket addr */

...
/* listenfd will be an endpoint for all requests to port
   on any IP address for this host */
if (bind(listenfd, (SA *)&serveraddr, sizeof(serveraddr)) < 0)
    return -1;
```

# Echo Server: `open_listenfd` `(listen)`

- `listen` indicates that this socket will accept connection (connect) requests from clients
- `LISTENQ` is constant indicating how many pending requests allowed

```
int listenfd; /* listening socket */

...
/* Make it a listening socket ready to accept connection requests */
if (listen(listenfd, LISTENQ) < 0)
    return -1;
return listenfd;
}
```

- We're finally ready to enter the main server loop that accepts and processes client connection requests.


# Echo Server: Main Loop

- The server loops endlessly, waiting for connection requests, then reading input from the client, and echoing the input back to the client.

```
main() {  
  
    /* create and configure the listening socket */  
  
    while(1) {  
        /* Accept(): wait for a connection request */  
        /* echo(): read and echo input lines from client til EOF */  
        /* Close(): close the connection */  
    }  
}
```



# Overview of the Sockets Interface



# Echo Server: accept


- `accept ( )` blocks waiting for a connection request

```
int listenfd; /* listening descriptor */
int connfd;   /* connected descriptor */
struct sockaddr_in clientaddr;
int clientlen;


clientlen = sizeof(clientaddr);
connfd = Accept(listenfd, (SA *)&clientaddr, &clientlen);
```

- `accept` returns a *connected descriptor* (`connfd`) with the same properties as the *listening descriptor* (`listenfd`)
  - Returns when the connection between client and server is created and ready for I/O transfers
  - All I/O with the client will be done via the connected socket
- `accept` also fills in client's IP address


# Echo Server: `accept` Illustrated



1. Server blocks in `accept`, waiting for connection request on listening descriptor `listenfd`



2. Client makes connection request by calling and blocking in `connect`



3. Server returns `connfd` from `accept`. Client returns from `connect`. Connection is now established between `clientfd` and `connfd`

# Connected vs. Listening Descriptors

- Listening descriptor
  - ▣ End point for client connection requests
  - ▣ Created once and exists for lifetime of the server
  
- Connected descriptor
  - ▣ End point of the connection between client and server
  - ▣ A new descriptor is created each time the server accepts a connection request from a client
  - ▣ Exists only as long as it takes to service client
  
- Why the distinction?
  - ▣ Allows for concurrent servers that can communicate over many client connections simultaneously
    - E.g., Each time we receive a new request, we fork a child to handle the request

# Echo Server: Identifying the Client

- The server can determine the domain name, IP address, and port of the client

```
struct hostent *hp; /* pointer to DNS host entry */
char *haddrp;      /* pointer to dotted decimal string */
unsigned short client_port;
hp = Gethostbyaddr((const char *)&clientaddr.sin_addr.s_addr,
                  sizeof(clientaddr.sin_addr.s_addr), AF_INET);
haddrp = inet_ntoa(clientaddr.sin_addr);
client_port = ntohs(clientaddr.sin_port);
printf("server connected to %s (%s), port %u\n",
       hp->h_name, haddrp, client_port);
```

# Echo Server: echo

- The server uses RIO to read and echo text lines until EOF (end-of-file) is encountered.
  - ▣ EOF notification caused by client calling `close(clientfd)`

```
void echo(int connfd)
{
    size_t n;
    char buf[MAXLINE];
    rio_t rio;

    Rio_readinitb(&rio, connfd);
    while((n = Rio_readlineb(&rio, buf, MAXLINE)) != 0) {
        upper_case(buf);
        Rio_writen(connfd, buf, n);
        printf("server received %d bytes\n", n);
    }
}
```

# Testing Servers Using `telnet`

- The `telnet` program is invaluable for testing servers that transmit ASCII strings over Internet connections
  - ▣ Our simple echo server
  - ▣ Web servers
  - ▣ Mail servers
- Usage:
  - ▣ `unix> telnet <host> <portnumber>`
  - ▣ Creates a connection with a server running on `<host>` and listening on port `<portnumber>`

# Testing the Echo Server With telnet

```
greatwhite> echoserver 15213
```

```
linux> telnet greatwhite.ics.cs.cmu.edu 15213
```

```
Trying 128.2.220.10...
```

```
Connected to greatwhite.ics.cs.cmu.edu.
```

```
Escape character is '^['.
```

```
hi there
```

```
HI THERE
```



## For More Information

- W. Richard Stevens, “Unix Network Programming: Networking APIs: Sockets and XTI”, Volume 1, Second Edition, Prentice Hall, 1998
  - ▣ THE network programming bible
- Unix Man Pages
  - ▣ Good for detailed information about specific functions
- Complete versions of the echo client and server are developed in the text
  - ▣ Updated versions linked to course website
  - ▣ Feel free to use this code in your assignments